

Matrix Distributed Processing 2.X

Massimo Di Pierro*

School of Computer Science, Telecommunications and Information Systems
DePaul University, 243 S. Wabash Av., Chicago, IL 60604 USA

April 18, 2001

Abstract

We present a set of programming tools (classes and functions written in C++ and based on Message Passing Interface) for fast development of generic parallel (and non-parallel) lattice simulations. They are collectively called MDP.

These programming tools include classes and algorithms for matrices, random number generators, distributed lattices (with arbitrary topology), fields and parallel iterations. No previous knowledge of MPI is required in order to use them.

Some applications in electromagnetism, electronics, condensed matter and lattice QCD are presented.

Keywords: lattice, parallel computing, numerical software

*Web page: www.phoenixcollective.org/mdp/

Program Summary

- *Title of the library:* MDP, Matrix Distributed Processing
- *Computer for which the library is designed:* Any computer or parallel computer, including clusters of workstations.
- *Software required in order to run MDP:* MPI (Message Passing Interface) is required to run MDP in parallel. MPI is not required to run MDP in single process mode.
- *Computer(s) on which the library has been tested.:* SUN SparcSTATION with Solaris, Linux PCs, cluster of PCs (connected by Ethernet and Myrinet). The code has also been tested partially on a Cray T3E.
- *Programming language used:* ANSI C++, to be compiled with

g++ (GCC) version 2.95.3

There are known problems with gcc 3.x.x, Intel and VC++ compilers.

- *Parallel applications provided together with the library as examples:* A program that solves electrostatic problems (`application1.cpp`). A program that computes total impedance in a net of resistors (`application2.cpp`). An Ising model simulation (`application3.cpp`). A complete Lattice QCD package (`FermiQCD`, not described here).

- *MDP web site:*

http://www.phoenixcollective.org/mdp/index_mdp.html

- *FermiQCD web site:*

<http://www.fermiqcd.net>

Contents

1	Introduction	5
1.1	Motivations	5
1.2	Content	6
2	Non Parallel Tools	10
2.1	Using <code>mdp_array</code>	11
2.2	Using <code>mdp_matrix</code>	12
2.3	class <code>mdp_random_generator</code>	16
2.4	class <code>mdp_jackboot</code>	18
3	Parallel Tools	21
3.1	What's MPI?	21
3.2	<code>mpi</code> on top of MPI	23
3.3	<code>mpi</code> and Wormholes	25
3.3.1	<code>get()</code> and <code>put()</code>	27
3.3.2	<code>put/get</code> matrices	29
3.3.3	Global communications ★	29
4	Introducing Lattices and Fields	32
4.1	class <code>mdp_lattice</code>	32
4.2	class <code>mdp_site</code>	35
4.3	class <code>mdp_field</code>	38
4.3.1	Local random generator	40
4.3.2	Accessing a lattice from a field	41
4.3.3	More on <code>update()</code>	41
4.3.4	<code>operator()</code> ★	42
4.3.5	A few derived fields	43
4.4	Console Input/Output	44
4.5	File Input/Output	45
4.6	The internal representation of a lattice ★	47
4.7	Partitioning and topology ★	52
4.7.1	Non periodic lattice	54
4.7.2	Hexagonal lattice	54
4.7.3	Antiperiodic boundary conditions	56
4.8	Memory optimization ★	57

5	Examples of parallel applications	57
5.1	Electrostatic potential for a given distribution of charges . . .	57
5.2	Total impedance in net of resistors	59
5.3	Ising model	62
5.4	Solid under stress	67
5.5	Lattice QCD	67
6	Timing and efficiency issues	70
A	Syntax of MDP	74

1 Introduction

MDP (Matrix Distributed Processing) is a collection of classes and functions written in C++ and based on Message Passing Interface. It was originally developed to implement QCD simulations on cluster of PCs using a natural matrix formalism.

It evolved in time and became a stand-alone general purpose package for fast development of any kind of parallel numerical simulations that involve the lattice discretization of a finite portion of space. A number of tricks and optimizations are implemented that make of MDP an efficient tool for developing, among others, fast and robust Lattice QCD applications.

1.1 Motivations

One of the biggest discoveries of modern physics is that long distance forces are mediated by fields and these fields are subject to local interactions. Local differential equations, in some field variables, seems to describe any sector of fundamental physics (from particle physics to gravitation, from thermodynamics to condensed matter).

Often the equations describing complex interacting systems cannot be solved exactly and require a numerical approach, i.e. discretize on a lattice the space on which the equations are defined but, the more complex is the system, the more computing power is required.

The good thing of having to deal with local equations is that one can partition the problem by partitioning the lattice (the space) on which the problem is defined, and only a minimal set of communications is required to pass information about the field variables at the boundary of these partitions (because first and second derivatives of the differential equations in the continuum correspond to nearest neighbor interaction terms on the lattice).

Parallel computers and clusters are the natural tool for implementing the numerical algorithms that can solve this kind of problems. Parallel computers with shared memory present, in fact, a bottle-neck in accessing the memory and, since the equations are local, there is no real need for a shared memory.

Clusters of PCs have many advantages over big parallel computers:

- They are made with commodity hardware therefore they are cheap.

- Their use is so wide spread that they come with state-of-the-art processors (almost 1GHz in speed) and a lot of memory.
- They are easy to design and upgrade.
- They can run open source operative systems such as Linux.
- They are “standard” in the sense that different manufacturers produce different processors and motherboards that are compatible with each other at both software and hardware level. This implies that one can upgrade/replace the hardware without necessary having to upgrade/replace the software.

Communication bandwidth within the cluster may be an issue but, even if communications increase with the number of PCs in the cluster, for the kind of problems we are considering, the ratio communication/computation actually decreases when the size of the problem increases! (exactly as the ratio surface/volume of a box decreases when the size of the box increases).

Said this, we believe that today the cost of developing, debugging and optimizing parallel software for large scale numerical lattice simulations (including the cost for computing power and human resources) remains a major issue to be addressed. A cost that has become comparable with the cost of running the simulations themselves.

MDP has been developed to solve this problem. Its use can considerably reduce that development cost since one can concentrate on the algorithm instead of focusing on communications and programming tricks to optimize efficiency.

For example MDP has been used here at Fermilab to develop, in a relatively short time, a parallel software for Lattice QCD simulations (see section 5.5).

1.2 Content

The main characteristics of MDP are:

- It includes `mdp_matrix`, a class to manipulate matrices of complex (`mdp_complex`) numbers. The assignment operator and the copy constructor have been overloaded to optimize the speed and the memory usage. In fact a `mdp_matrix` object contains a dynamically allocated pointer to the location of memory where the matrix elements

are stored. The matrix elements are automatically copied only when it is inevitable, otherwise their memory address is copied. In particular, the temporary `mdp_matrix` object returned by a function stores its elements in the same location of memory used by the function that will take that object as argument. This class is used to implement an interface to the fields.

- All the calls to MPI are done through a wrapper object, `mpi`, which takes care of generating communication tags and computing the length of the objects to be sent/received by the parallel processes.
- The most important classes in MDP are

```
mdp_lattice  
mdp_site  
mdp_field
```

They enable the user to define lattices of any size, shape and topology and associate any kind of structure to the sites. Other properties of these classes are: automatic communications between the parallel processes (when necessary), parallel I/O for the field variables and functions to move on the lattice.

- For each user defined lattice, one random generator per site is automatically created and is a member variable of `mdp_lattice`. This guarantees that any simulation will give the same results whether it runs on a single PC or in parallel, independently of the way the lattice is partitioned. The random generator itself can be seen as a field and it can be saved and restored and any point.
- Each lattice (and each field defined on that lattice) can be partitioned arbitrarily among the different processes (or automatically if the user does not bother to do it) and the constructor of `mdp_lattice`, for each process, takes care of allocating the memory for storing the neighbor sites that have to be passed by a neighbor process. The code has been optimized in order to minimize the need for communications and the low level optimization tricks are completely hidden to the high level programmer.

- The same code can run on single process on one PC or multiprocess on a cluster of PCs (or on a supercomputer supporting MPI), without any modification. The results are independent from the number of processes¹.

As an example of what one can do with MDP we present a parallel program that creates a field of 5×5 matrices on a $4 \times 4 \times 4$ lattice (distributed on 4 processors) and sets each matrix (for each site) equal to the inverse of the exponential of a random $SU(5)$ matrix (just to make things a little complicated!). Then it averages each matrix with its 6 nearest neighbors.

```

01: // Program: example01.cpp
02: #include "mdp.h"
03:
04: int main(int argc, char **argv) {
05:     mdp.open_wormholes(argc,argv);
06:     int          ndim=3;
07:     int          mybox[]={4,4,4};
08:     mdp_lattice  mylattice(ndim, mybox);
09:     mdp_matrix_field F(mylattice,5,5);
10:     site        x(mylattice);
11:
12:     forallsites(x)
13:         F(x)=inv(exp(mylattice.random(x).SU(5)));
14:     F.update();
15:     forallsites(x) {
16:         F(x)=1.0/7.0*(F(x)+F(x+0)+F(x-0)+F(x+1)+F(x-1)+F(x+2)+F(x-2));
17:         cout << "F(x) for x=(" << x(0) << ", " << x(1) << ")\n";
18:         cout << F(x) << endl;
19:     }
20:     mdp.close_wormholes();
21:     return 0;
22: }
```

Note that each line has a meaning and none of them is superfluous:

¹Apart for rounding errors that usually are compiler and hardware dependent.

- Line 2 includes the libraries of MDP
- Line 5 establishes the communications among the processes
- Line 6 sets a variable containing the dimensions of the lattice
- Line 7 fills an array containing the size of the lattice
- Line 8 defines the lattice, `mylattice`
- Line 9 defines a field of 5×5 matrices `F` on `mylattice` (`mdp_matrix_field` is built-in field based on `mdp_field`).
- Line 10 defines a site variable `x` on `mylattice`
- Lines 12,13 initialize each field variable `F(x)` with the inverse of the exponential of a random $SU(5)$ (generated using the local random generator of site `x`).
- Line 14 performs all the communication to update boundary sites, i.e. sites that are shared by different processes. They are determined automatically and optimally when the lattice is declared.
- Lines 15,16 average `F(x)` with its 6 next neighbor sites
- Lines 17,18 print out the results
- Line 20 closes all the communications

The source code of MDP is contained the `mdp_*.h` files. Only part of the library is described in this tutorial.

Some sections of this manual are marked with a star. This indicates that an important technical detail is discussed in the section but it can be skipped in a first reading.

A reader interested only in parallelization issues can jump directly to section 4, and go back to sections 2 and 3 later.

The last page of this tutorial provides information to request the source code of MDP.

2 Non Parallel Tools

The most important classes declared by MDP are:

- `mdp_real`. It is equivalent to `float` unless the user defines the global constant `USE_DOUBLE_PRECISION`. In this case `mdp_real` stands for `double`.
- `mdp_complex`. Declared as `complex<mdp_real>`. The imaginary unit is implemented as a global constant `I=mdp_complex(0,1)`.
- `mdp_array<object,n>`. It is a container for `n`-dimensional arrays of objects (for examples arrays of `mdp_matrix`). `mdp_arrays` can be resized any time, moreover they can be passed to (and returned by) functions.
- `mdp_matrix`. An object belonging to this class is effectively a matrix of `mdp_complex` numbers and it may have arbitrary dimensions.
- `mdp_random_generator`. This class contains the random number generator and member functions to generate random `float` numbers with uniform distribution, Gaussian distribution or any user defined distribution. It can also generate random $SU(n)$ matrices. Different `mdp_random_generator` objects can be declared and initialized with different seeds.
- `mdp_jackboot`. It is a class to store sets of data to be used for computing the average of any user defined function acting on the data. It includes member functions for computing Jackknife and Bootstrap errors on this average.

Moreover the following constants are declared

```
#define and      &&
#define or      ||
#define Pi      3.14159265359
#define I       mdp_complex(0,1)
#define CHECK_ALL
#define TRUE    1
#define FALSE   0
```

If the definition of `CHECK_ALL` is removed the code will be faster but some checks will be skipped resulting a lesser safe code.

2.1 Using `mdp_array`

`mdp_arrays` are declared using templates. For example the command

```
mdp_array<float,4> myarray(2,3,6,5);
```

declares a $2 \times 3 \times 6 \times 5$ array of `float` called `myarray`. The first argument of the template specifies the type of object (i.e. `float`) and the second argument is the number of dimensions of the array (i.e. 4). An array of this kind can be resized at any time. For example

```
myarray.dimension(4,3,2,2);
```

transforms `myarray` in a new $4 \times 3 \times 2 \times 2$ array of `float`. A `mdp_array` can have up to 10 dimensions. The array elements can be accessed by reference using `operator()` in the natural way

```
myarray(i,j,k,l)
```

where `i,j,k,l` are integers. The total number of dimensions (in this example 4) is contained in the member variable

```
myarray.ndim
```

Two useful member functions are

```
myarray.size()
```

that returns the total number of elements of the array (as `long`) and

```
myarray.size(n)
```

that returns the size of the of the dimension `n` of the array (as `long`). The argument `n` must be an integer in the range `[0,ndim-1]`.

Here is an example of a function that returns an `mdp_array`.

```
mdp_array<float,2> f(float x) {  
    mdp_array<float,2> a(2,2);  
    a(0,0)=x;  
    return a;  
}
```

Here is an example of a function that takes an `mdp_array` as argument

```
float g(mdp_array<float,2> a, int i, int j) {  
    return a(i,j);  
}
```

2.2 Using mdp_matrix

A `mdp_matrix` object, say `M`, is very much like `mdp_array<mdp_complex,2>`. A `mdp_matrix` can be declared either by specifying its size or not

```
mdp_matrix M(r,c);    // r rows times c columns
mdp_matrix M;         // a general matrix
```

Even if the size of a matrix has been declared it can be changed anywhere with the command

```
M.dimension(r,c);    // r rows times c columns
```

Any `mdp_matrix` is automatically resized, if necessary, when a value is assigned to it. The following lines

```
mdp_matrix M(5,7), A(8,8);
M=A;
cout << M.rowmax() << "," << M.colmax() << endl;
```

prints 8,8, The member functions `rowmax()` and `colmax()` return respectively the number of rows and columns of a `mdp_matrix`.

The element (i,j) of a matrix `M` can be accessed with the natural syntax

`M(i,j)`

where `i,j` are integers. Moreover the class contains functions to perform standard operations among matrices:

`+, -, *, /, +=, -=, *=, /=, inv, det, exp, sin, cos, log,`
`transpose, hermitiam, minor, identity,...`

As an example a program to compute

$$\left[\exp \begin{pmatrix} 2 & 3 \\ 4 & 5i \end{pmatrix} \right]^{-1} \quad (1)$$

looks like this:

```
// Program: example02.cpp
#include "mdp.h"

int main() {
    mdp_matrix a(2,2);
    a(0,0)=2; a(0,1)=3;
    a(1,0)=4; a(1,1)=5*I;
    cout << inv(exp(a)) << endl;
    return 0;
}
```

Here is one more example of how to use the class `mdp_matrix`:

```
// Program: example03.cpp
#include "mdp.h"

mdp_matrix cube(Matrix X) {
    mdp_matrix Y;
    Y=X*X*X;
    return Y;
}

int main() {
    mdp_matrix A,B;
    A=Random.SU(3);
    B=cube(A)*exp(A)+inv(A);
    cout << A << endl;
    cout << B << endl;
    return 0;
}
```

This code prints on the screen a random $SU(3)$ matrix A and $B = A^3 e^A + A$. Some example statements are listed in tab. 1. Note that the command

```
A=mul_left(B,C);
```

is equivalent but faster than

```
A=C*transpose(B);
```

Example	C++ with MDP_Lib2.h
$A \in M_{r \times c}(\mathbf{C})$	<code>A.dimension(r,c)</code>
A_{ij}	<code>A(i,j)</code>
$A = B + C - D$	<code>A=B+C-D</code>
$A^{(ij)} = B^{(ik)}C^{(kj)}$	<code>A=B*C</code>
$A^{(ij)} = B^{(jk)}C^{(ik)}$	<code>A=mul_left(B,C)</code>
$A = aB + C$	<code>A=a*B+C</code>
$A = a\mathbf{1} + B - b\mathbf{1}$	<code>A=a+B-b</code>
$A = B^T C^{-1}$	<code>A=transpose(B)*inv(C)</code>
$A = B^\dagger \exp(iC)$	<code>A=hermitian(B)*exp(I*C)</code>
$A = \cos(B) + i \sin(B) * C$	<code>A=cos(B)+I*sin(B)*C</code>
$a = \text{real}(\text{tr}(B^{-1}C))$	<code>a=real(trace(inv(B)*C))</code>
$a = \det(B)\det(B^{-1})$	<code>a=det(B)*det(inv(B))</code>

Table 1: Examples of typical instructions acting on `mdp_matrix` objects. `A,B,C,D` are assumed to be declared as `mdp_matrix`; `r,c` as `int`; `a,b` may be any kind of number.

because it does not involve allocating memory for the transposed of B.

As one more example of a powerful application of `mdp_array` here is a program that defines a multidimensional `mdp_array` of `mdp_matrix` objects and pass it to a function.

```
// Program: example04.cpp
#include "mdp.h"

mdp_array<mdp_matrix,3> initialize() {
    mdp_array<mdp_matrix,3> d(20,20,20);
    int i,j,k;
    for(i=0; i<20; i++)
        for(j=0; j<20; j++)
            for(k=0; k<20; k++) {
                d(i,j,k).dimension(2,2);
                d(i,j,k)(0,0)=k;
                d(i,j,k)(0,1)=i;
                d(i,j,k)(1,0)=j;
                d(i,j,k)(1,1)=k;
            }
    return(d);
}

mdp_array<mdp_matrix,3> f(mdp_array<mdp_matrix,3>& c) {
    mdp_array<mdp_matrix,3> d(c.size(0),c.size(1),c.size(2));
    int i,j,k;
    for(i=0; i<c.size(0); i++)
        for(j=0; j<c.size(1); j++)
            for(k=0; k<c.size(2); k++)
                d(i,j,k)=sin(c(i,j,k));
    return(d);
}

int main() {
    mdp_array<mdp_matrix,3> a, b;
    a=initialize();
    b=f(a);
}
```

```

    int i=1, j=2, k=3;
    cout << a(i,j,k) << endl;
    cout << b(i,j,k) << endl;
    return 0;
}

```

a, b, c, d are 3D $20 \times 20 \times 20$ arrays of 2×2 matrices. The program prints

```

[[ 3.000+0.000i  1.000+0.000i ]
 [ 2.000+0.000i  3.000+0.000i ]]
[[ 0.022+0.000i  -0.691+0.000i ]
 [ -1.383+0.000i  0.022+0.000i ]]

```

Note that the instruction

```
d(i,j,k)=sin(c(i,j,k));
```

computes the `sin()` of the `mdp_matrix` `c(i,j,k)`.

2.3 class `mdp_random_generator`

The random generator implemented in `mdp_random_generator` is the Marsaglia random number generator described in ref.[6] (the same generator is also used by the UKQCD collaboration in many large scale numerical simulations). It presents the nice feature that it can be initialized with a single long random number and its correlation time is relatively large compared with the standard random generator `rand()` of C++. To define a random generator, say `myrand`, and use the integer `seed` as seed one should simply pass `seed` to the constructor:

```
mdp_random_generator myrand(seed);
```

For simple non-parallel applications one only needs one random generator. For this reason one `mdp_random_generator` object called `Random` is automatically created and initialized with seed 0.

`mdp_random_generator` contains some member variables for the seeds and four member functions:

- `plain()`; It returns a random `float` number in the interval $[0,1)$ with uniform distribution.
- `gaussian()`; It returns a random `float` number x generated with a Gaussian probability $P(x) = \exp(-x^2/2)$.
- `distribution(float (*P)(float, void*), void* a)`; It returns a random number x in the interval $[0,1)$ generated with a Gaussian probability $P(x, a)$, where `a` is any set of parameters pointed by `void *a` (the second argument passed to `distribution`). The distribution function P should be normalized so that $0 \leq \min P(x) \leq 1$ and $\max P(x) = 1$.
- `mdp_random_generator::SU(int n)`; it returns a random `mdp_matrix` in the group $SU(n)$ or $U(1)$ if the argument is $n = 1$.

Here is a simple program that creates one random number generator, generates a set of 1000 random Gaussian numbers and counts how many of them are in the range $[n/2, (n+1)/2)$ for $n = 0 \dots 9$

```
// Program: example05.cpp
#include "mdp.h"

int main() {
    mdp_random_generator random;
    int i,n,bin[10];
    float x;
    for(n=0; n<10; n++) bin[n]=0;
    for(i=0; i<1000; i++) {
        x=random.gaussian();
        for(n=0; n<10; n++)
            if((x>=0.5*n) && (x<0.5*(n+1))) bin[n]++;
    }
    for(n=0; n<10; n++)
        cout << "bin[" << n << "] = " << bin[n] << endl;
    return 0;
}
```

Here is a program that computes the average of the sum of two sets of 1000 numbers, `a` and `b`, where `a` is generated with probability $P(a) = \exp(-(a - \bar{a})^2/(2\sigma))$ and `b` with probability $Q(b) = \sin(\pi b)$.

```

// Program: example06.cpp
#include "mdp.h"

float Q(float x, void *a) {
    return sin(Pi*x);
}

int main() {
    mdp_random_generator random;
    int i,N=100;
    float a,b,average=0, sigma=0.3, a_bar=1;
    for(i=0; i<N; i++) {
        a=(sigma*random.gaussian()+a_bar);
        b=random.distribution(Q);
        average+=a+b;
        cout << "average=" << average/(i+1) << endl;
    }
    return 0;
}

```

For large N the output asymptotically approaches $a_bar+0.5=1.5$.

The algorithm for $SU(2)$ is based on map between $O(3)$ and $SU(2)$, realized by

$$\{\hat{a}, \alpha\} \rightarrow \exp(i\alpha\hat{a} \cdot \sigma) = \cos(\alpha) + i\hat{a} \cdot \sigma \sin(\alpha) \quad (2)$$

where $(\sigma^1, \sigma^2, \sigma^3)$ is a vector of Pauli matrices, $\hat{a} \in S^2$ is a uniform vector on the sphere and $\alpha \in [0, \pi)$ is a uniform rotation angle around that direction.

A random $SU(n)$ is generated using the well-known Cabibbo-Marinari iteration [3] of the algorithm for $SU(2)$.

2.4 class mdp_jackboot

Suppose one has n sets of m different measured float quantities $x[i][j]$ (where $i = 0 \dots n$ and $j = 0 \dots m$) and one wants to compute the average over i , of a function $F(x[i])$. For example

```

float x[n][m]={...}
float F(float *x) {

```

```

    return x[1]/x[0];
};
float result=0;
for(i=0; i<n; i++)
    result+=F(x[i])/n;
};

```

Then one may ask: what is the error on the result? In general there are two algorithms to estimate this error: Jackknife and Bootstrap [4]. They are both implemented as member functions of the class `mdp_jackboot`. They work for any arbitrary function F .

A `mdp_jackboot` object, let's call it `jb`, is a sort of container for the data. After it has been filled it can be asked to return the mean of $F()$ and its errors. Here is an example of how it works:

```

mdp_jackboot jb(n,m);
jb.f=F;
for(i=0; i<n; i++)
    for(j=0; j<m; j++)
        jb(i,j)=x[i][j];
cout << "Result          = " << jb.mean() << endl;
cout << "Jackknife error = " << jb.j_err() << endl;
cout << "Bootstrap error = " << jb.b_err(100) << endl;

```

Note that

- The constructor of the class `mdp_jackboot` takes two arguments: the first is the number of configuration; the second is the number of the quantities measured on each configuration.
- `jb.f` is the pointer to the function used in the analysis.
- `jb.mean()` returns the mean.
- `jb.j_err()` returns the Jackknife error.
- `jb.b_err()` returns the Bootstrap error. It takes as argument the number of Bootstrap samples. The default value is 200.

It is possible to declare arrays of `mdp_jackboot` objects, but it is rarely necessary. It is simpler to declare different functions and repeat the analysis using the same `mdp_jackboot` object assigning the pointer `mdp_jackboot::f` to each of the functions at the time.

As another example consider the following program. It generates an array of 100 $SU(6)$ matrices. For each matrix it computes trace and determinant, and returns the average of the ratio between the real part of the trace and the real part of the determinant (with its Jackknife and Bootstrap errors), and the average of the product of the real part of the trace and the real part of the determinant (with its Jackknife and Bootstrap errors)

```
// Program: example07.cpp
#include "mdp.h"

const int n=100;

float f1(float *x, void *a) { return x[0]/x[1]; }

float f2(float *x, void *a) { return x[0]*x[1]; }

int main() {
    mdp_matrix A;
    mdp_jackboot jb(n,2);
    int i;
    for(i=0; i<n; i++) {
        A=Random.SU(6)+Random.gaussian();
        jb(i,0)=real(det(inv(A)));
        jb(i,1)=real(det(A));
    }
    jb.f=f1;
    cout << "Result x[0]/x[1] = " << jb.mean() << endl;
    cout << "Jackknife error = " << jb.j_err() << endl;
    cout << "Bootstrap error = " << jb.b_err(100) << endl;
    jb.f=f2;
    cout << "Result x[0]*x[1] = " << jb.mean() << endl;
    cout << "Jackknife error = " << jb.j_err() << endl;
    cout << "Bootstrap error = " << jb.b_err(100) << endl;
}
```

```

    return 0;
}

```

Note that any user defined function used by `mdp_jackboot` (`f1` and `f2` in the example) has to take two arguments: an array of `float` (containing a set of measurements done on a single configuration) and a pointer to `void`. This is useful to pass some extra data to the function. The extra data must be passed to `mdp_jackboot` by assigning the member variable

```
void* mdp_jackboot::handle;
```

to it.

`mdp_jackboot` has one more member function `plain(int i)`, where the call

```

/* define mdp_jackboot jb(...) */
/* assign a value to the integer i */
jb.plain(i)

```

is completely equivalent to

```

/* define mdp\_jackboot jb(...)          */
/* assign a value to the integer i */
float f(float *x, void *a) {
    return x[i];
};
jb.f=f

```

Using `plain` saves one from defining trivial functions for `mdp_jackboot`.

3 Parallel Tools

3.1 What's MPI?

With a parallel program we mean a job constituted by many programs running in parallel (eventually on different processors) to achieve a global goal. Each of the running programs, part of the same job, is called a process. Different processes can communicate to each other by exchanging messages (in

MPI) or by accessing each-other memory (in computers with shared memory). The latter case is hardware dependent and will not be taken in consideration. With the term “partitioning” we will refer to the way different variables (in particular the field variables defined on a lattice) are distributed among the different processes.

The most common and portable parallel protocol is Message Passing Interface (MPI). It is implemented on a number of different platforms, for example Unix, Linux, Solaris, Windows NT and Cray T3D/E.

When writing a program using MPI, one essentially writes one program for each processor (Multiple Instructions Multiple Data) and MPI provides the functions for sending/receiving information among them.

We present here an example of how MPI works, even if a detailed knowledge of MPI is not required to understand the rest of this tutorial.

Suppose one wants to compute $(5 * 7) + (4 * 8)$ in parallel using two processes. A typical MPI program to do it is the following:

```
#include "stdio.h"
#include "mpi.h"
int main(int argc, char **argv) {
    int ME, Nproc;
    MPI_Status status;
    MPI_Init(&argc, &argv);
    MPI_Comm_rank(MPI_COMM_WORLD, &ME);
    MPI_Comm_size(MPI_COMM_WORLD, &Nproc);
    if(ME==1) {
        int b;
        b=4*8;
        MPI_Send(&b,1,MPI_INT,0,45,MPI_COMM_WORLD);
    };
    if(ME==0) {
        int a,b;
        a=5*7;
        MPI_Recv(&b,1,MPI_INT,1,45,MPI_COMM_WORLD,&status);
        printf("%i\n", a+b);
    };
    MPI_Finalize();
};
```

`ME` is a variable that contains the unique identification number of each running process and `Nproc` is the total number of processes running the same code. Since different processes have different values of `ME` they execute different parts of the program (`if(ME==n) ...`).

In the example $(5 * 7)$ is computed by process 0 while, at the same time $(4 * 8)$ is computed by process 1. Then process 1 sends (`MPI_Send`) its partial result to process 0, which receives it (`MPI_Recv`) and prints out the sum².

MPI instructions can be classified in

- initialization functions (like `MPI_Init`, `MPI_Comm_rank`, `MPI_Comm_size`)
- one to one communications (like `MPI_Send`, `MPI_Recv`)
- collective communication (where for example one variable can be broadcasted by one process to all the others)

It is not a purpose of this manual to explain how MPI works therefore we refer to [7].

3.2 mpi on top of MPI

MPI is not Object Oriented³ and it is so general that MPI function calls usually require a lot of arguments. This makes it very easy to do mistakes when programming with MDP.

For this reason the class `mpi_wormhole_class`, a wrapper to MPI, has been created. It has only one global representative object called `mpi`. We believe that `mpi` is easier to use than MPI in the context of lattice simulation.

`mpi` is not intended as a substitute or an extension to MPI, in fact only a subset of the features of MPI are implemented in `mpi` and occasionally one may need explicit MPI calls. Its main purpose is to build an intermediate programming level between MPI and `mdp_lattice/mdp_field`, which constitute the very essence of MDP. There are some advantages of having this intermediate level, for example:

²Needless to say that in such a simple program the parallelization is completely useless because the message passing takes more time than the computation of the whole expression!

³There is an Object Oriented version of MPI, called OOMPI, but we decided not to use it because it not as standard as MPI is.

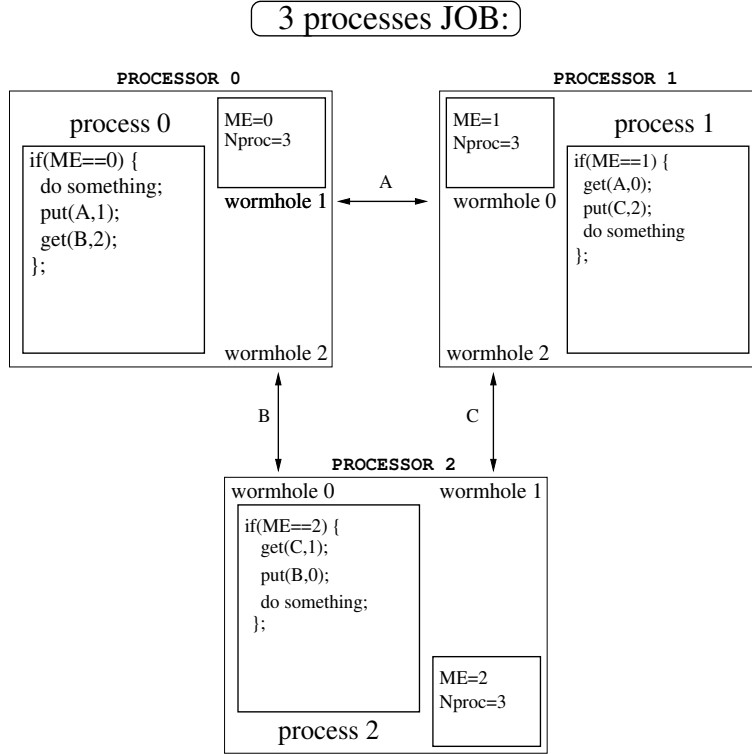


Figure 1: Example of a 3 processes jobs. ME and Nproc are the global variables that characterize the different processes running (eventually) on different processors. The term wormhole refers to the virtual connection between each couple of processes.

- It allows one to compile the code without linking with **MPI**. In this case **mpi** substitutes some dummy functions to the **MPI** calls. This allows one to compile and run any parallel program on a single processor computer even if **MPI** is not installed, and it is also important for testing and debugging the code.
- In case one wants to run **MDP** on a parallel machine that does not support **MPI** but does support a different protocol, one does not have to rewrite everything but only minor modifications to the class **mpi_wormhole_class** will be necessary.

3.3 **mpi** and Wormholes

From now on we will use here the word “wormhole”, in its intuitive sense, to represent the virtual connections between the different processes. One can think of each process to have one wormhole connecting it with each of the other processes. Each wormhole has a number, the number of the process connected to the other end. Process X can send something to process Y just putting something into its wormhole n.Y. Process Y has to get that something from its wormhole n.X. **mpi** is the global object that allows the user to access these wormholes (put/get stuff to/from them). A schematic representation of a 3 processes job is shown in fig. 1.

Using **mpi**, instead of **MPI**, program above becomes:

```
// Program: example08.cpp
#include "mdp.h"

int main(int argc, char **argv) {
    mdp.open_wormholes(argc,argv);
    if(mdp.nproc()==2) {
        if(ME==1) {
            int b;
            b=4*8;
            mdp.put(b,0);
        }
        if(ME==0) {
            int a,b;
```

```

        a=5*7;
        mdp.get(b,1);
        mpi << a+b << endl;
    }
} else {
    mpi << "Sorry, this only runs on 2 processes\n";
}
mdp.close_wormholes();
}

```

The most important things to notice are:

1. To compile with MPI use the `mpiCC` compiler and the compiler option `-DPARALLEL`. When it is omitted some dummy functions are substituted to the MPI calls and communications are skipped. This is useful for testing the behavior of a particular process in a multiprocess job without actually going parallel.
2. `Nproc` is the total number of processes. It is a macro defined as `mdp.nproc()`
3. Many standard headers (`stdio.h`, `math.h`, `complex.h`, `iostream.h`) are included by `mdp.h`.
4. `ME` is a global macro defined as `mdp.id()`. It represents the identification number of the running process. Its value is assigned by the call to `mdp.open_wormholes` when running in `PARALLEL`, otherwise it is zero.
5. All of the initialization functions have been replaced by a single one:

```
void mdp_wormhole_class::open_wormholes(int&,char**);
```

It is compulsory even if the program does not run in parallel. It takes as input the same arguments of `main()`⁴.

⁴This is due to the fact that when running in parallel (using the `runmpi` command) a script submits the same program to the different processors with different command parameters, which are passed as arguments to `main()`. This is how a process knows its own identification number: `ME`.

6. `MPI_Send` and `MPI_Recv` have been replaced by `mdp.put` and `mdp.get` respectively. They check automatically the size and type of their arguments.
7. At the end of the program one has to `close_wormholes`. Actually if one forgets about this the destructor of `mpi` does take care of it.

The programmer is forced to use the names `Nproc` and `ME` for the total number of processes and the process identification number respectively. They are keywords for MDP.

3.3.1 `get()` and `put()`

The functions `put/get` have the following syntax

```
template <class T> mdp_wormhole_class::put(T&,int);
template <class T> mdp_wormhole_class::put(T&,int,mpi_request&);
template <class T> mdp_wormhole_class::get(T&,int)
```

The first argument (passed by reference) is the object to be sent (`put`) or where to store the received one (`get`). The second argument is the destination (`put`) or the source (`get`). Forget about `mpi_request` for the moment, it is optional.

All the variables related with message passing are members of `mpi`. In this way variables that are not explicitly necessary remain hidden inside it (for example the communicator and the communication tag).

The functions `put` and `get` that replaced `MPI_Send` and `MPI_Recv` are Object Oriented in the sense that they can be used to send any kind of object⁵. The size and the tag of the communication are automatically computed and do not have to be specified by the programmer. For example with the same syntax a `mdp_complex` number can be `put/get` instead of an integer. The only price paid for this is that a process, X, cannot send more than one object to another process, Y, until the first object has been received. This is not such a price because in the kind of computations we are interested in, we want to minimize the number of `put/get` and we want to pack as much data as possible in each single communication, instead of making many small ones.

⁵Assuming it does not contain a pointer to dynamically allocated memory.

To avoid confusion, if the third argument of `put()` is not specified this function will not return until the send process has been completed (synchronous send)⁶. Sometimes one wants the `put()` command to return even if the send is not completed, so that the process is free to do more computations, and eventually check at a later stage if the previous `put()` has completed its task (asynchronous send). One typical case is when the object to be send was dynamically allocated and it cannot be deallocated until the communication is concluded. The right way to implement this send is writing a code like the following:

```
/* mystruct is any built in variable
   or user defined structure */
mystruct *myobj=new mystruct;
mpi_request request;
put(*myobj,destination,request);
/* do something else */
mdp.wait(request);
delete myobj;
```

The `wait()` member function of `mpi` stops the process until the `put()` command has terminated its task. Its only argument is the object belonging to the class `mpi_request` that was passed to the corresponding `put()`. The function `wait()` can also take a pointer to an array of requests and wait for all of them. The general definition of this function is

```
void mdp_wormhole_class::wait(mpi_request&);
void mdp_wormhole_class::wait(mpi_request*, int);
```

the second argument is the length of the array.

It is also possible to put/get arrays of objects by using the overloaded functions

```
template <class T> mdp_wormhole_class::put(*T,long,int)
template <class T> mdp_wormhole_class::put(*T,long,int,mpi_request&)
template <class T> mdp_wormhole_class::get(*T,long,int)
```

⁶Actually in `mdp.put` an synchronous send is internally implemented using `MPI_Isend` + `MPI_Wait`. This is because `MPI_Send` is ambiguous in some cases: the same program can have different behavior depending on the size of the objects that are passed. Our implementation avoids any ambiguity.

where the first argument is a pointer to the first element of the array, the second is the length, the third is the destination/source.

3.3.2 put/get matrices

We have said that with the put/get commands it is not possible to pass objects that contain a pointer to dynamically allocated memory and this is the case for a mdp_matrix object. This does not mean that one cannot put/get a mdp_matrix. The correct way to send a mdp_matrix is

```
mdp\_matrix A;  
/* do something with A */  
mdp.put(A.address(),A.size(),destination);
```

The correct way to receive a mdp_matrix is

```
mdp\_matrix B;  
/* dimension B as A */  
mdp.get(B.address(),B.size(),destination);
```

It is extremely important to stress that when one puts/gets a matrix one only sends/receives its elements. Therefore (in the last example) B must have the same dimensions of A at the moment `get` is called, otherwise one gets wrong results.

3.3.3 Global communications ★

As an example of global communication we will consider the following code

```
// Program: example09.cpp  
#include "mdp.h"  
  
int main(int argc, char **argv) {  
    mdp.open_wormholes(argc,argv);  
    int a;  
    if(mdp.nproc()==2) {  
        if(ME==1) {  
            a=4*8;  
            mdp.add(a);  
        }  
    }  
}
```

```

    }
    if(ME==0) {
        a=5*7;
        mdp.add(a);
        cout << "a=" << a << endl;
    }
} else {
    mpi << "Sorry, this only runs on 2 processes\n";
}
mdp.close_wormholes();
}

```

The member function:

```
void mdp_wormhole_class::add(float&);
```

sums the first argument passed by all the processes. In this way each process knows the result of the sum operation.

To sum an array of float the corresponding sum for arrays of float is implemented as:

```
void mdp_wormhole_class::add(float*,long);
```

(the second argument is the length of the array). The same function `add` also works for double, long, int, `mdp_complex` and `mdp_matrix` and arrays of these types. To add `mdp_complex` numbers:

```
mdp_complex x;
mdp.add(x);
```

To add `mdp_matrix` objects

```
mdp_matrix A;
mdp.add(A);
```

Some more member functions of `mpi` for collective communications are:

- `barrier()` which sets a barrier and all processes stop at that point until all the processes reach the same point.
- `abort()` which forces all the processes to abort.

- `template<class T> broadcast(T a, int p)` which broadcast the object `a` (belonging to an arbitrary class `T`) of process `p` to all other processes.
- `template<class T> broadcast(T *a, long n, int p)` which broadcast the array `a` of objects `T` of process `p` to all other processes. `n` here is the number of elements of the array.

Here is an example of how to use these collective communications

```
// Program: example13.cpp
#include "mdp.h"

int main(int argc, char **argv) {
    mdp.open_wormholes(argc,argv);
    int i,j;
    for(i=0; i<5; i++) {
        if(ME==0) j=i;
        else      j=0;
        if(i%2==0) mdp.barrier();
        mdp.broadcast(j,0);
        cout << "I am process " << ME
        << ", i=" << i
        << ", j=" << j << endl;
    }
    mdp.close_wormholes();
}
```

This programs prints

```
I am process 0, i=0, j=0
I am process 0, i=1, j=1
I am process 0, i=2, j=2
I am process 0, i=3, j=3
I am process 0, i=4, j=4
I am process 1, i=0, j=0
I am process 1, i=1, j=1
I am process 1, i=2, j=2
I am process 1, i=3, j=3
I am process 1, i=4, j=4
```

4 Introducing Lattices and Fields

4.1 class `mdp_lattice`

By a lattice we mean a generic set of points and a set of functions to move from one point to another. In MDP a lattice is implemented as any subset of a regular n-dimensional grid with a given (arbitrary) topology.

A lattice of such a kind is implemented as an object belonging to the class `mdp_lattice` using the command:

```
mdp_lattice mylattice(ndim,mybox,mypartitioning,  
                      mytopology,seed,next_next,local_random);
```

or the command

```
mdp_lattice mylattice(ndim,ndir,mybox,mypartitioning,  
                      mytopology,seed,next_next,local_random);
```

where:

- `mylattice` is a user defined variable that will contain the actual lattice structure.
- `ndim` is the dimension of the basic regular grid.
- `ndir` is the number of directions. It can be omitted and, by default, it is assumed to be equal to `ndim`.
- `mybox` is a user defined `ndim`-ensional array of `int` that contains the size of the basic regular grid.
- `mypartitioning` is a user defined function that for each site of the basic grid returns the number of the process that stores it (or it returns `NOWHERE` if the point has to be excluded from the lattice). If `mypartitioning` is omitted all sites of the original grid (specified by `mybox`) will be part of the final lattice, and they will be equally distributed between the processes according with the value of coordinate 0 of each point⁷.

⁷For example, consider a lattice of size 8 in the 0 direction distributed on 4 processes. Process 0 will contain sites coordinate 0 equal to 0 and 1. Process 1 will contain sites coordinate 0 equal to 2 and 3. An so on.

- `mytopolgy` is a user defined function that for each site of the final (remaining) grid, and for each dimension in the `mydim`-ensional space, returns the coordinates of the neighbor sites in the up and down direction. If `mytopology` is omitted the lattice will, by default, have the topology of an `ndim`-ensional torus, i.e. the same of the basic grid plus periodic boundary conditions).
- `seed` is an integer that is used to construct a site dependent seed to initialize the random generators associated to each lattice sites. If omitted it is assumed to be 0.
- `next_next` is an integer that may have only three values (1,2 or 3). It essentially fixes the thickness of the boundary between different processes. If omitted the default value is 1. It's use will not be discussed in this tutorial.
- `local_random` is a boolean value, true by default. It specify is a local random generator (true) or a global random generator (false) are to be used for the lattice. The second (false) option results in a faster program and less RAM usage but it has issues on parallel programs.
- Note that each lattice stores topology information in RAM memory. Sometime this may be inconvenient if the lattice is be excessively big. In this case the use of the compiler option `-DMDP_NO_LG` ensures that topology information is saved on disk and not stored in RAM. For small lattices this may slow down the program. For big lattices it only slows down the initialization phase.

For the moment we will restrict ourselves to the study of lattices with the basic (torus) topology, postponing the rules concerning the specifications for `mypartitioning` and `mytopolgy`.

Here is a short program to create a 2D 8^2 lattice (but no field yet associated to it) running in parallel on 4 processes.

```
// Program: example11.cpp
#include "mdp.h"

int main(int argc, char **argv) {
    mdp.open_wormholes(argc,argv);
```

```

int mybox[]={8,8};
generic_lattice mylattice(2,mybox);
mdp.close_wormholes();
}

```

Assuming the program is running on 4 processors, when it runs the lattice will be automatically subdivided into 4 sublattices of size $\{2,8\}$. `mylattice` contains information about the sites and how they are distributed among the different processes. Moreover the different processes automatically communicate to each other this information so that each of them knows which local sites will have to be passed to the neighbor processes. All the communications to establish the topology are done only once and the information are inherited by all the fields defined on `mylattice`. The program produces the following output

```

PROCESS 0 STARTING
PROCESS 2 STARTING
PROCESS 1 STARTING
PROCESS 3 STARTING
=====
Starting [ Matrix Distributed Processing ] ...
Created by Massimo Di Pierro (mdp@metacryption.com)
=====
Going parallel ... YES

Initializing a mdp_lattice...
Communicating...
Initializing random per site...
Done. Let's begin to work!
=====
Process 0 stats: 10%; 25%; 94%.
Process 1 stats: 9%; 26%; 99%.
Process 2 stats: 8%; 27%; 92%.
Process 3 stats: 10%; 25%; 91%.
=====
PROCESS 1 ENDING AFTER 0.010 sec.
PROCESS 2 ENDING AFTER 0.008 sec.
PROCESS 3 ENDING AFTER 0.011 sec.

```

PROCESS 0 ENDING AFTER 0.012 sec.

On Linux machines the use of the compiler option `-DLINUX` ensures that program output includes usage statistics. For example the line

```
Process 0 stats: 10%; 25%; 34%.
```

shows that processor 0 is using 10%. These statistics may help to identify slow processors and bugs in programs.

Two other member functions of `mdp_lattice` are

```
long mdp_lattice::global_volume();  
long mdp_lattice::local_volume();
```

The first returns the size of the total lattice volume (i.e. the total number of sites), and the second returns the size of the portion of volume stored by the calling process.

The member function

```
long mdp_lattice::size();
```

returns the total volume of the box containing the lattice and

```
long mdp_lattice::size(int mu);
```

returns the size, in direction μ , of the box containing the lattice.

4.2 class `mdp_site`

To access a site one can define a `mdp_site` variable using the syntax

```
mdp_site x(my_lattice);
```

This tells the compiler that `x` is a variable that will contain the coordinate of a site and a reference to the lattice `my_lattice` on which the point is defined.

Let's look at a modification of the preceding code that creates a lattice and, for each site `x`, prints out the identification number of the process that stores it:

```
// Program: example12.cpp
#include "mdp.h"

int main(int argc, char **argv) {
    mdp.open_wormholes(argc,argv);
    int mybox[]={10,10};
    generic_lattice mylattice(2,mybox);
    site x(mylattice);
    forallsites(x)
        cout << "Site=(" << x(0) << "," << x(1)
        << ") is stored by " << ME << endl;
    mdp.close_wormholes();
}
```

Note that

`x(mu)`

returns the coordinate `mu` of `x`.

This simple program allows one to print a map of the lattice. The output of this code depends on the number of processors on which it is running⁸. The loop command

```
forallsites(x) /* do something with x */
```

spans first all the even sites and then all the odd sites.

It is possible to set the value of a `site` variable `x` to a particular site of given coordinates `((3,7)` for example) using the command

```
x.set(3,7);
```

but one must be careful and execute this statement only on the process which contains the site `(3,7)` or, eventually, in a process that contains a neighbor of this site. In this case, in fact, the process will also contains a copy of the original site therefore it will be able to access the location. By default each process will store its own sites plus a copy of next-neighbor sites in each direction and a copy of next-next neighbor sites (moving in two different directions)⁹. They will be referred to as boundary sites. Boundary sites are

⁸The buffers of the different processors are copied to the standard output at a random time partially scrambling the output of the different processes. For this reason it is in general a good rule to print only from process `ME==0`.

⁹This is required, for example, by the clover term in QCD.

different for different processes.

It is possible to check if a site is a local site in the following way:

```
x.set(3,7)
if(x.is_in()) /* then do something */
```

The member function `is_in()` returns `TRUE` if the site is local. A better way to code it would be

```
if(on_which_process(3,7)==ME) {
    x.set(3,7);
    /* do something */
}
```

Other member functions of `mdp_site` are

- `is_in_boundary()` returns `TRUE` if a site is in the boundary of that process.
- `is_here()` returns `TRUE` if the site `is_in()` or `is_in_boundary()`.
- `is_equal(int,int,int...)` returns `TRUE` if the site is equal to the site specified by the coordinates that are arguments of `is_equal()` (`x0,x1,x2,...`). **Accessing a site that `!is_here()` crashes the program.**

To move from one site to another is quite simple:

```
x=x+mu; // moves x up in direction mu (integer) of one step
x=x-mu; // moves x down in direction mu (integer) of one step
```

Note that `mu` is an integer and in general

$$x + \mu \neq x \neq x - \mu \quad (3)$$

even when `mu` is zero.

Here is a test program to explore the topology:

```
// Program: example13.cpp
#include "mdp.h"
```

```

int main(int argc, char **argv) {
    mdp.open_wormholes(argc,argv);
    int mybox[]={10,10};
    mdp_lattice mylattice(2,mybox);
    mdp_site x(mylattice);
    int mu=0;
    if(ME==0) {
        x.set(0,0);
        do {
            cout << "x=(" << x(0) << "," << x(1) << ")\n";
            if(x.is_in_boundary()) error("I found the boundary");
            x=x+mu;
        } while(true);
    }
    mdp.close_wormholes();
}

```

This program stops after 5 iterations if `Nproc==2`, after 4 if `Nproc==3`, after 3 if `Nproc==4`, after 2 if `Nproc==5` and so on. It will never stop on a single process job (`Nproc==1`) because of periodic boundary conditions. Since all sites with the same `x(0)`, by default, are stored in the same process, if one moves in direction `mu=1` the job will never stop despite the number of processes.

4.3 class mdp_field

A `mdp_field` is the simplest field one can define on a given `mdp_lattice`.

Suppose, for example, one wants to associate a given structure, `mystruct`, to each site of `mylattice` (in the example a structure that contains just a `float`) and initialize the field variable to zero:

```

// Program: example14.cpp
#include "mdp.h"

struct mystruct {
    float value; /* or any other structure */
};

```

```

int main(int argc, char **argv) {
    mdp.open_wormholes(argc,argv);
    int mybox[]={10,10};
    mdp_lattice      mylattice(2,mybox);
    mdp_field<mystruct> myfield(mylattice);
    mdp_site          x(mylattice);
    forallsites(x)
        myfield(x).value=0;
    myfield.update();
    mdp.close_wormholes();
}

```

The command

```
mdp_field<mystruct> myfield(mylattice);
```

defines `myfield`, a field of `mystruct` on the sites of `mylattice`. The command `myfield(x)` returns by reference the structure associated to the site `x`

The function

```
void mdp_lattice::update(int,int,int)
```

is the most important of all. For the moment we will restrict to that case when it is called with no arguments. Its task is to perform all the communications necessary in order for each process to get an updated copy of the boundary sites stored by a neighbor process.

After the loop all sites that are local (`forallsites`) are initialized, but sites that are in boundary (since they are copies of sites initialized on a different process) still contain “random” numbers. One way to initialize also the sites in the boundary and avoid some time consuming communications is replacing the lines

```

forallsites(x)
    myfield(x).value=0;
mylattice.update();

```

with

```
forallsitesandcopies(x) myfield(x).value=0;
```

(Note that the command `forallsitesandcopies` only works for a local expression that does not call the local random generator).

It is also possible to loop on all the sites of a given parity, for example EVEN (or ODD), in a slow way

```
int parity=EVEN;
forallsites(x)
    if(x.parity()==parity)
        /* then do something */
```

or in a fast way

```
int parity=EVEN;
forallsitesofparity(x,parity)
    /* do something */
```

The second expression is faster because sites with the same parity are stored contiguously in memory.

4.3.1 Local random generator

Once a `mdp_lattice` is declared one random generator per site is automatically created and initialized with a function of the site coordinates. The random generator (and its member functions) of site `x` of `mylattice` can be accessed (by reference) with the command

```
mdp_random_generator& mdp_lattice::random(mdp_site)
```

Program `example14.cpp` can be modified so that `myfield` is initialized, for example, with a gaussian random number:

```
// Program: example15.cpp
#include "mdp.h"

struct mystruct {
    float value; /* or any other structure */
};

int main(int argc, char **argv) {
```



```

mdp.open_wormholes(argc,argv);
int mybox[]={10,10};
mdp_lattice      mylattice(2,mybox);
mdp_field<mystruct> myfield(mylattice);
mdp_site x(mylattice);
forallsites(x)
    myfield(x).value=mylattice.random(x).gaussian();
myfield.update();
mdp.close_wormholes();
}

```

4.3.2 Accessing a lattice from a field

It is always possible to access a `mdp_lattice` from its field using `lattice()`, a member function of `mdp_field` that returns by reference the `mdp_lattice` on which the field is defined.

For example in the preceding program one could substitute

```
myfield(x).value=mylattice.random(x).gaussian();
```

with

```
myfield(x).value=myfield.lattice().random(x).gaussian();
```

and get the same result. This is useful because one can avoid passing a `mdp_lattice` object to a function of a `mdp_field`.

4.3.3 More on update()

In some cases it may be useful to have a field that contains `n mystruct` variables at each site.

A field of this type can be defined with the command

```
mdp_field<mystruct> myfield(mylattice,n);
```

Consider, for example, the simple case of a tensor field $T_{\mu}^{ij}(x)$, where $i, j \in \{0..5\}$, $\mu \in \{0..3\}$. It can be defined with

```
mdp_field<mdp_complex[4][5][5]> T(mylattice);
```

and its elements can be accessed with the natural expression:

```
T(x) [mu] [i] [j]
```

where μ, i, j are integers.

Alternatively the tensor T_{μ}^{ij} can be defined with the command

```
mdp_field<mdp_complex[5][5]> T(mylattice,4);
```

and its elements can be accessed with the expression:

```
T(x,mu) [i] [j]
```

Sometimes the second choice is better for the following reason: In many algorithms one often needs to update all the boundary site variables having a given a parity and a given μ . One wants to do the update in one single communications. The `update()` function allows one to do exactly this:

```
int parity=EVEN, int mu=3;
T.update(parity,mu);
```

4.3.4 operator() ★

Some care is required when declaring a `mdp_field`. The structure associated to the sites cannot contain a pointer to dynamically allocated memory, therefore it cannot contain a `mdp_matrix` (or a `tt mdp_array`) object. This is not a limitation since the structure associated to the sites can be a multidimensional array and one can always map it into a `mdp_matrix` by creating a new field, inheriting `mdp_field` and redesigning `operator()`.

In this subsection we explain how to use `mdp_matrix` to build and interface to the field variables.

Let's go back to the example of a tensor field $T_{\mu}^{ij}(x)$. One can define it as

```
class mytensor: public mdp_lattice<mdp_complex[5][5]> {};
mytensor T(mylattice,4);
```

Instead of accessing its elements as `mdp_complex` numbers one may want to access the two-dimensional array as a `mdp_matrix` object. This is done by overloading the `operator()` of the class `mdp_field`:

```
mdp_matrix myfield::operator() (mdp_site x, int mu) {
    return mdp_matrix(address(x,mu),5,5);
};
```

where `address()` is a member function of `mdp_lattice` that returns the location of memory where the field variables at `(x,mu)` are stored. After this definitions one can simply access the field using the expression:

```
T(x,mu)
```

that now returns a `mdp_matrix`.

Now `mytensor T` looks like a field of 5×5 matrices and its elements can be used as discussed in chapter 2. For example one can initialize the tensor using random $SU(5)$ matrices generated independently by the local random generators associated to each site

```
forallsites(x)
    for(mu=0; mu<4; mu++)
        T(x,mu)=T.lattice().random(x).SU(5);
```

or print a particular element

```
x.set(3,5);
if(x.is_in()) print(T(x,2));
```

4.3.5 A few derived fields

Following the directives of the last subsections four more basic fields are implemented in MDP.

- `mdp_matrix_field`: a field of `mdp_matrix`;
- `mdp_nmatrix_field`: a field of `n mdp_matrix`;
- `mdp_vector_field`: a field of vectors (a vector is seen as a `mdp_matrix` with one single column);
- `mdp_nvector_field`: a field of `n vectors`;

To explain their usage we present here a few lines of code that define a field of 3×4 matrices, **A**, a field of 4-vectors, **u**, and a field of 3-vectors, **v**; then for each site compute

$$v(x) = A(x)u(x) \quad (4)$$

This can be implemented as

```
mdp_matrix_field A(mylattice,3,4);
mdp_vector_field u(mylattice,4);
mdp_vector_field v(mylattice,3);
/* assign values to the fields */
forallsites(x) v(x)=A(x)*u(x);
```

and its generalization to

$$v_i(x) = A_i(x)u_i(x) \quad (5)$$

for *i* in the range [0,N-1]

```
int N=10; // user defined variable
mdp_nmatrix_field A(mylattice,N,3,4);
mdp_nvector_field u(mylattice,N,4);
mdp_nvector_field v(mylattice,N,3);
/* assign values to the fields */
forallsites(x) for(i=0; i<N; i++) v(x,i)=A(x,i)*u(x,i);
```

4.4 Console Input/Output

When running in parallel one should be careful when printing to the console. Typically one may want to specify that only one process prints to the console and not all of them. The built-in stream **mdp**, also called **mpi** serves this purpose. It can be used anywhere in the program by only if **ME==0** it connects to **cout**.

It is possible to enable/disable the stream **mdp** for any processor by turning true/false the flag **mdp.print**.

For example in:

```
if(ME==0 || ME==3) mdp.print=true;
mdp << "Hello World\n";
```

only processes 0 and 3 print “Hello World”.

4.5 File Input/Output

The class `mdp_field` contains two I/O member functions:

```
void mdp_lattice::save(char[],int,int);  
void mdp_lattice::load(char[],int,int);
```

They take three arguments

- The `filename` of the file to/from which to save/load the data
- The identification number of the master process that is will physically perform to the I/O
- The size (in sites) of the buffer used by the master process to communicate with the other processes.

The last two arguments are optional and by default the master process is process 0 and the buffer size is 1024 sites. The identification number of the master process has to correspond to the process running on the processor that is physically connected with the I/O device where the file `filename` is. The size of the buffer will not affect the result of the I/O operation but may affect the speed (which should increase with the size of the buffer). If the size is too big there may be an “out of memory” problem (that usually results in obscure error messages when compiling with MPI). Note that the I/O functions save/load have to be called by all processes (not just the master one). In fact all processes which contain sites of the lattice are involved in the I/O by exchanging informations with the master.

As an example of I/O we consider the case of a `save()` operation. Each of the processes arranges the local field variables to be saved into packets and sends the packets to the process that performs the I/O (we assume it is process 0). Process `i` arranges its site into packets ordered according with the global parametrization and sends one packet at the time to process 0. In this way process 0 only receives one packet at the time for each of the processes and already finds the sites stored in the correct order so that it does not need to perform a seek. Process 0 saves the sites according to the global parametrization (which is independent from the lattice partitioning over the parallel processes). Its only task is to sweep all the sites and, for

each of them, pick the corresponding field variables from the packet sent by the process that stored it.

We believe this is the most efficient way to implement a parallel I/O for field variables. As a particular case, assuming process 0 has enough memory, one could set the size of the buffer equal to the maximum number of sites stored by a process. In this case the parallel I/O would be done performing only one communication per process.

As an example of efficiency: a field as large as 100MB can be read and distributed over 8 500MHz PentiumIII PCs in parallel (connected with Ethernet) in a few of seconds.

The save/load function are inherited by every field based on `mdp_field`. For example:

```
/* define mylattice */
class myfield: public mdp_field<mdp\_complex[7][3]> {};
myfield F(mylattice);
/* do something */
F.save("test.dat");
F.load("test.dat");
```

If a field contains member variables other than the structure at the sites, these member variables are ignored by save/load. The only variables that are saved together with the field are:

- the number of dimensions
- the size of the box containing the lattice
- the total number of sites belonging to the lattice
- the size in bytes of the structure at each lattice site
- a code to detect the endianness of the computer that saved the data
- the time and date when the file was created

We also provide a small program, `inspect.cpp`, that can open a file and extract this information.

It is always possible to add any other information to the bottom of the file (in binary or ASCII format) without compromising its integrity.




Figure 2: Example of a 6×6 grid. The points are labelled by their coordinates.

4.6 The internal representation of a lattice ★

A lattice is a collection of sites and, for sake of simplicity, from now on we will refer to the lattice defined by

```
int ndim=2;
int mybox[]={6,6};
int mypartitioning(int x[], int ndim, int nx[]) {
    if(x[0]<3) return 0;
    if(x[1]<4) return 1;
    return 2;
};
mdp_lattice mylattice(ndim,mybox,mypartitioning);
```

This lattice is represented in fig.2 and partitioned as in fig.5.

To speed up the communication process MDP uses different parametrizations for labelling the sites:

- the global coordinates (as shown in fig.2)
- a global index (as shown in fig.3)




Figure 3: Example of a 6×6 grid. The points are labelled by their global index.

- a local index (as shown in fig.8 from the point of view of process 0)

and it provides functions to convert one parametrization into another.

The global parametrization is quite natural. For example: the point $\mathbf{x}[] = \{3, 2\}$ has a global index given by $\mathbf{x}[1] * \text{mybox}[0] + \mathbf{x}[0]$; The two global parametrizations have a one to one correspondence. The local parametrization is more complicated and it represents the order in which each process stores the sites in the memory.

When the lattice is defined the constructor of the class `mdp_lattice` spans all the sites using the global coordinates (fig.2) and it assigns them the global index (fig.3) and a parity (fig.4). Then each process checks which sites are local sites (fig.5). After the topology is assigned to the lattice (fig.6), each process identifies which sites are neighbors of the local sites (fig.7 for process 0). At this point each process has all the information needed to build a local parametrization for the sites. All even sites belonging to process 0 are labeled with a progressive number; then all the odd sites from process 0 are labelled in progression. The same procedure is repeated by each process for each process (including itself). The result (for the case in the example) is shown in fig.8 (from the point of view of process 0).




Figure 4: Example of a 6×6 grid. The points are labelled by their global index. Points with even (white) and odd (black) parity are distinguished.

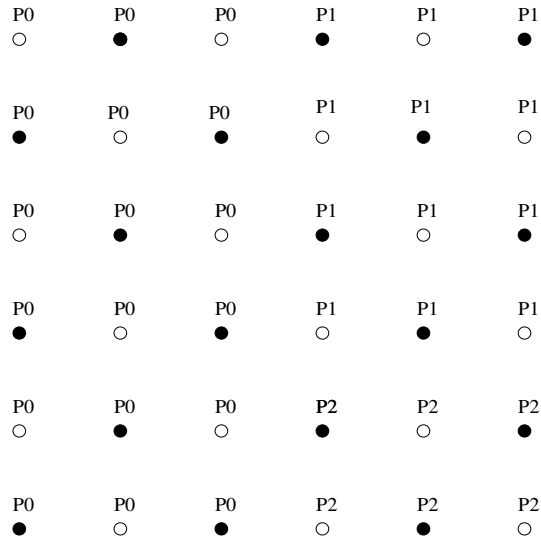


Figure 5: Example of a 6×6 grid distributed on 3 processors (according with some user defined partitioning). Each point is labelled by the identification number of the process that stores it.




Figure 6: Example of a 6×6 grid endowed with the default (torus) topology. When the topology is assigned to the grid, it is promoted to the rank of lattice.

When a field is associated to the lattice the field variables are stored by each process according with the local parametrization. This guarantees that:

- Boundary sites that have to be copied from the same process are stored continuously in the memory. This minimizes the need for communications.
- Boundary sites copied from the same process (and the local sites as well) of given parity are also stored continuously in memory. This speeds up loops on sites of given parity.
- When a process X copies a neighbor site from process Y it does not need to allocate a buffer because the data can be received directly in the memory location where it is normally stored.

Given a site variable `x` on `mylattice` one can ask for the `mu` coordinate using the syntax

`x(mu)`

or the global index

local ○	local ●	local ○	from P1 ●	○	from P1 ●
local ●	local ○	local ●	from P1 ○	●	from P1 ○
local ○	local ●	local ○	from P1 ●	○	from P1 ●
local ●	local ○	local ●	from P1 ○	●	from P1 ○
local ○	local ●	local ○	from P2 ●	○	from P2 ●
local ●	local ○	local ●	from P2 ○	●	from P2 ○

Figure 7: Example of a 6×6 lattice form the point of view of process 0. Process 0 knows which sites are local, which non-local sites are neighbors of the local sites and from which process to copy them.

0 ○	9 ●	1 ○	22 ●	?	23 ●
10 ●	2 ○	11 ●	18 ○	?	19 ○
3 ○	12 ●	4 ○	24 ●	?	25 ●
13 ●	5 ○	14 ●	20 ○	?	21 ○
6 ○	15 ●	7 ○	28 ●	?	29 ●
16 ●	8 ○	17 ●	26 ○	?	27 ○

Figure 8: Example of a 6×6 lattice form the point of view of process 0. Process zero assigns a local parametrization to the local sites and to the neighbor sites of the local points. It cannot access the rest of the sites.

```

    x.global_index()
or the local index
    x.local_index()
Moreover one can initialize a site variable specifying the coordinates
    x.set(3,2);
the global index
    x.set_global(26);
or the local index
    x.set_local(7); /* assuming ME is 0 */
If x and y are two site variables in the same code defined on the same lattice,
but using different partitioning. One can assign y to x using the assignment
operator
    x=y;
This is equivalent to
    x.set_global(y.global_index());

```

4.7 Partitioning and topology ★

One of the most important characteristics of MDP is that one is not limited to a box-like lattice (even if it has to be a subset of the points in a box) and the sites of the lattice can be associated to an arbitrary topology. Before stating with weird stuff here is the definition of the default function for the lattice partitioning:

```

template<int dim>
int default_partitioning(int *x, int ndim, int *nx) {
    float tpp1=nx[dim]/NPROC;
    int tpp2=(int) tpp1;
    if(tpp1-tpp2>0) tpp2+=1;
    return x[dim]/tpp2;
};

```

Any partitioning function takes three arguments:

- An array `x[ndim]` containing the `ndim` coordinates of a site
- The dimensions, `ndim`, of the box containing the lattice sites
- An array, `nx[ndim]`, containing the size of box in each dimension

It returns the identification number of the process that stores the site corresponding to the coordinates in `x[]`¹⁰.

One can define a different partitioning just by defining another function that takes the same arguments as the default one and pass it as third argument to the constructor of `mdp_lattice`. If the new partitioning function, for a particular input site, returns a number outside the range 0 to `NPROC-1`, that site is excluded from the lattice. To this scope the macro `NOWHERE` is defined. If a site has to be excluded from the lattice the partition function should return `NOWHERE`. If one defines a partitioning function that excludes some points from the lattice one is forced to create a user defined topology, otherwise one ends up with sites which are neighbors of sites that are `NOWHERE` and the program fails.

The default torus topology is defined as

```
void torus_topology(int mu, int *x_dw, int *x, int *x_up,
                  int ndim, int *nx) {
    for(int nu=0; nu<ndim; nu++) if(nu==mu) {
        x_dw[mu]=(x[mu]-1+nx[mu]) % nx[mu];
        x_up[mu]=(x[mu]+1) % nx[mu];
    } else x_up[nu]=x_dw[nu]=x[nu];
};
```

Any topology function takes six arguments, they are:

- A direction `mu` in the range 0 to `ndim-1`
- A first array of coordinates, `x_dw[]`
- A second array of coordinates, `x[]`
- A third array of coordinates, `x_up[]`
- The dimensions, `ndim`, of the box containing the lattice sites
- An array, `nx[ndim]`, containing the size of box in each dimension

¹⁰that at this level `site` variables are not used. In fact the partitioning function (as well the topology function) is called before the lattice is defined, and it is used to define it. A site variable can be defined only after a lattice is defined.

The coordinates in `x[]` are taken as input and the topology function fills the the arrays `x_dw[]` and `x_up[]` with the coordinates of the neighbor points when moving from `x[]` one step in direction `mu` down and up respectively.

The topology function will only be called for those sites `x[]` that belong to the lattice. It has to fill `x_dw[]` and `x_up[]` with the coordinates of points that have not been excluded from the lattice.

4.7.1 Non periodic lattice

As an example one can create a lattice with a true physical boundary by saying, for example, that moving up (or down) in a particular direction from a particular subset of sites one remains where one is¹¹. In the simple case of a finite box one could define the topology:

```
void box_topology(int mu, int *x_dw, int *x, int *x_up,
                 int ndim, int *nx) {
    torus_topology(mu,x_dw,x,x_up,ndim,nx);
    if(x[mu]==0) x_dw[mu]=x[mu];
    if(x[mu]==nx[mu]-1) x_up[mu]=x[mu];
};
```

4.7.2 Hexagonal lattice

Another example of a lattice with a weird topology is the hexagonal grid shown in fig.9. The trick to do it is having a 3D lattice flat in the third direction, put NOWHERE the centers of the hexagons and define the proper topology for the remaining points. It can be done by the following code

```
int exagonal_partitioning(int x[], int ndim, int nx[]) {
    if((x[0]+x[1])%3==0) return NOWHERE;
    else return default_partitioning<0>(x,ndim,nx);
};

void exagonal_topology(int mu, int *x_dw, int *x, int *x_up,
                      int ndim, int *nx) {
    if((x[0]+x[1])%3==1) {
```

¹¹This kind of lattice can be used, for example, to study a 3D solid under stress and study its deformations.




Figure 9: Example of a 6×6 grid to be used to build an hexagonal lattice (top) and the metric associated to this grid (bottom). the points at center of the hexagons are excluded from the lattice by the partitioning function.

```

switch(mu) {
case 0: x_up[0]=(x[0]+1)%nx[0]; x_up[1]=x[1]; break;
case 1: x_up[1]=(x[1]+1)%nx[1]; x_up[0]=x[0]; break;
case 3: x_up[0]=(x[0]-1+nx[0])%nx[0];
      x_up[1]=(x[1]-1+nx[1])%nx[1]; break;
};
x_up[2]=x[2]; x_dw[0]=x[0];
x_dw[1]=x[1]; x_dw[2]=x[2];
};
if((x[0]+x[1])%3==2) {
switch(mu) {
case 0: x_dw[0]=(x[0]-1+nx[0])%nx[0]; x_dw[1]=x[1]; break;
case 1: x_dw[1]=(x[1]-1+nx[0])%nx[1]; x_dw[0]=x[0]; break;
case 3: x_dw[0]=(x[0]+1)%nx[0]; x_dw[1]=(x[1]+1)%nx[1]; break;
};
x_dw[2]=x[2]; x_up[0]=x[0];
x_up[1]=x[1]; x_up[2]=x[2];
};
};

int n=2;
int mybox[]={3*n, 3*n, 1};

mdp_lattice exagonal(3,mybox,exagonal_partitioning,
                    exagonal_topology);

```

One can observe that we need a flat third dimension to allow for three possible direction. For this reason MDP allows to define a lattice with a number of directions different from the number of dimensions. To avoid entering in a number of technical details we will avoid discussing this possibility here. What one can do with such a lattice is another matter and will not be discussed. From now on we will concentrate on more regular lattices which tend to have a more intuitive interpretation.

4.7.3 Antiperiodic boundary conditions

There are two possible meaning for antiperiodic boundary conditions: 1) antiperiodic boundary conditions for the lattice (i.e. a moebius topology).

2) antiperiodic boundary conditions for a field defined on the lattice.

The first case can be handled in the same way as any other topology. The second case has nothing to do with the lattice topology, but is a property of the field and one has to take care of it in the definition of `operator()` for the field in question.

4.8 Memory optimization ★

Actually the class `mdp_field` has more properties of those described. For example each field can be deallocated at any time if memory is needed:

```
struct mystruct { /* any structure */ };
mdp_lattice mylattice1(mydim1,mybox1);
mdp_field<mystruct> myfield(mylattice1);
/* use myfield */
myfield.deallocate_memory();
mdp_lattice mylattice2(mydim2,mybox2);
myfield.allocate_field(mylattice2);
```

Note that one cannot reallocate a field using a different structure at the site. This would be very confusing.

5 Examples of parallel applications

We will present here a few example of parallel application that use MDP. To keep things simple we will not make use of fields of matrices.

5.1 Electrostatic potential for a given distribution of charges

We consider here the vacuum space inside a cubic metal box connected to ground and containing a given (static) distribution of charge. We want to determine the electrostatic potential in the box.

The vacuum is implemented as a finite 20^3 lattice. Two fields, a charge density $q(x)$ and a potential $u(x)$, are defined on it. In the continuum the potential is obtained by solving the Gauss law equation

$$\nabla^2 u(x) = q(x) \tag{6}$$

It discretized form reads

$$\sum_{i=0}^2 \left[u(x + \hat{i}) - 2u(x) + u(x - \hat{i}) \right] = q(x) \quad (7)$$

which can be solved in $u(x)$

$$u(x) = \frac{1}{6} \left[q(x) + \sum_{i=0}^2 u(x + \hat{i}) + u(x - \hat{i}) \right] \quad (8)$$

Therefore the static potential solution is obtained by iterating eq.(8) on the vacuum until convergence.

As initial condition we assume that only two charges are present in the box:

$$q(x) = 3\delta(x - A) - 5\delta(x - B) \quad (9)$$

where $A = (3, 3, 3)$ and $B = (17, 17, 17)$. Moreover since the box has a finite extension the `box_topology` will be used.

Here is the parallel program based on MDP.

```
// Program: application1.cpp
#include "mdp.h"

int main(int argc, char **argv) {
    mdp.open_wormholes(argc, argv);
    int mybox[]={20,20,20};
    mdp_lattice vacuum(3, mybox,
                      default_partitioning<0>,
                      box_topology);
    mdp_field<float> u(vacuum);
    mdp_field<float> q(vacuum);
    mdp_site x(vacuum);
    mdp_site A(vacuum);
    mdp_site B(vacuum);
    A.set(3,3,3);
    B.set(17,17,17);
    float precision, old_u;
    forallsitesandcopies(x) {
```

```

    u(x)=0;
    if(x==A)      q(x)=3;
    else if(x==B) q(x)=-5;
    else          q(x)=0;
}
do {
    precision=0;
    forallsites(x) if((x(0)>0) && (x(0)<mybox[0]-1) &&
                      (x(1)>0) && (x(1)<mybox[1]-1) &&
                      (x(2)>0) && (x(2)<mybox[2]-1)) {
        old_u=u(x);
        u(x)=(q(x)+u(x+0)+u(x-0)+u(x+1)+u(x-1)+u(x+2)+u(x-2))/6;
        precision+=pow(u(x)-old_u,2);
    }
    u.update();
    mdp.add(precision);
} while (sqrt(precision)>0.0001);
u.save("potential.dat");
mdp.close_wormholes();
}

```

The saved values can be used to produce a density plot representing the potential in the volume.

5.2 Total impedance in net of resistors

Another problem we want to solve is that of determining the total resistance between two arbitrary points (A and B) on a semi-conducting finite cylindrical surface. The cylinder is obtained starting from a torus topology and cutting the torus in one direction (say 0). The total resistance is probed by connecting the terminals of a current generator (J) to the points A and B and evaluating the potential at points A and B using Ohm's law

$$R_{AB} = \frac{u(A) - u(B)}{J} \quad (10)$$

The surface of the semiconductor is implemented as a periodic net of resistances so that each site has four resistances connected to it, fig. 10. For




Figure 10: Example of a cylindrical net of resistors.

simplicity we assume that all the resistances are the same (i.e the material is homogeneous). A different choice would be equally possible.

The potential $u(x)$ is computed using the Kirchoff law at each vertex

$$\sum_{i=0,1} R_i(x - \hat{i}) \left[u(x - \hat{i}) - u(x) \right] + R_i(x) \left[u(x + \hat{i}) - u(x) \right] + J = 0 \quad (11)$$

that, solved in $u(x)$, leads to

$$u(x) = \frac{\sum_{i=0,1} R_i(x - \hat{i})u(x - \hat{i}) + R_i(x)u(x + \hat{i}) + J}{\sum_{i=0,1} R_i(x - \hat{i}) + R_i(x)} \quad (12)$$

This equation is iterated on each site until convergence.

Here is the parallel program based on MDP.

```
// Program: application2.cpp
#include "mdp.h"

void open_cylinder(int mu, int *x_dw, int *x, int *x_up,
                  int ndim, int *nx) {
```

```

    torus_topology(mu,x_dw,x,x_up,ndim,nx);
    if((mu==0) && (x[0]==0))      x_dw[0]=x[0];
    if((mu==0) && (x[0]==nx[0]-1)) x_up[0]=x[0];
}
float resistance(int x0, int x1, int mu) {
    return (Pi/100*x0);
}
int main(int argc, char **argv) {
    mdp.open_wormholes(argc,argv);
    int mybox[]={100,20};
    mdp_lattice cylinder(2,mybox,default_partitioning<0>,
                        open_cylinder);
    mdp_field<float> u(cylinder);
    mdp_field<float> r(cylinder,2);
    mdp_site x(cylinder);
    mdp_site A(cylinder);
    mdp_site B(cylinder);
    float precision, old_u;
    float c, J=1, deltaJ, deltaR, Rtot, local_Rtot;
    A.set(15,7);
    B.set(62,3);
    forallsitesandcopies(x) {
        u(x)=0;
        r(x,0)=resistance(x(0),x(1),0);
        r(x,1)=resistance(x(0),x(1),1);
    }
    do {
        precision=0;
        forallsites(x) {
            old_u=u(x);
            if(x==A) { c+=J; cout << ME << endl; };
            if(x==B) c-=J;
            deltaJ=c; deltaR=0;
            /* the next two lines take care of the finite cylinder */
            if(x+0!=x) { deltaJ+=u(x+0)*r(x,0);    deltaR+=r(x,0); };
            if(x-0!=x) { deltaJ+=u(x-0)*r(x-0,0); deltaR+=r(x-0,0); };
            deltaJ+=u(x+1)*r(x,1);    deltaR+=r(x,1);

```




Figure 11: Example of a 2D spin system.

```

        deltaJ+=u(x-1)*r(x-1,1); deltaR+=r(x-1,1);
        u(x)=deltaJ/deltaR;
        precision+=pow(u(x)-old_u,2);
    }
    u.update();
    mdp.add(precision);
} while (sqrt(precision)>0.00001);
Rtot=0;
if(A.is_in()) Rtot+=u(A)/J;
if(B.is_in()) Rtot-=u(B)/J;
mdp.add(Rtot);
mdp << "R_A-R_B=" << Rtot << endl;
mdp.close_wormholes();
}

```

5.3 Ising model

One more full application we wish to consider is the study of the total magnetization of a spin system under a magnetic field (this is also known as the Ising model) as function of the temperature. As usual we define a lattice that represents the geometry of the spin system, and associate to the sites a field $s(x) \in \{-1, +1\}$. This is schematically represented in fig. 11. The

Hamiltonian of this system is

$$H = -\beta \sum_{x,y} s(x)\Delta(x,y)s(y) - \kappa \sum_x s(x)[h_{tot} + h(x)] \quad (13)$$

where h_{tot} is the average external magnetic field, $h(x)$ is a local magnetic field, $\beta = 1/T$ is the inverse temperature and κ is a constant typical of the material (in principle one could promote κ to be a field for non-homogeneous materials. $\Delta(x,y)$ is 1 if $|x - y| = 1$, 0 otherwise. In the example we consider a two dimensional spin system but the extension to arbitrary dimensions is trivial: just change `mybox`. Moreover we implement a simple `montecarlo_multihit` algorithm [8] to minimize the action in presence of thermal fluctuations.

Here is the parallel program based on MDP .

```
// Program: application3.cpp
#include "mdp.h"

// //////////////////////////////////////
class scalar_field: public mdp_field<float> {
public:
    int ndim,nc;
    scalar_field(mdp_lattice &a) {
        allocate_field(a, 1);
        ndim=a.ndim;
    }
    float &operator() (mdp_site x) {
        return *(address(x));
    }
};

// //////////////////////////////////////
class ising_field: public mdp_field<int> {
public:
    int ndim,nc;
    float beta, kappa, magnetic_field;
    ising_field(mdp_lattice &a) {
        allocate_field(a, 1);
        ndim=a.ndim;
```

```

}
int &operator() (mdp_site x) {
    return *(address(x));
}
friend void set_cold(ising_field &S) {
    mdp_site x(S.lattice());
    forallsites(x) S(x)=1;
}

friend void montecarlo_multihit(ising_field &S, scalar_field &H,
int n_iter=1, int n_hits=1) {
    int iter, parity, hit, new_spin, mu;
    float delta_action;
    mdp_site x(S.lattice());
    for(iter=0; iter<n_iter; iter++) {
        for(parity=0; parity<=1; parity++)
            forallsitesofparity(x,parity) {
                for(hit=0; hit<n_hits; hit++) {
                    delta_action=S.kappa*(H(x)+S.magnetic_field);
                    for(mu=0; mu<S.ndim; mu++)
                        delta_action-=S(x+mu)+S(x-mu);
                    new_spin=(S.lattice().random(x).plain()>0.5)?1:-1;
                    delta_action*=S.beta*(new_spin-S(x));
                    if(delta_action<S.lattice().random(x).plain())
                        S(x)=new_spin;
                }
            }
        S.update(parity);
    }
}

friend float average_spin(ising_field &S) {
    float res=0;
    mdp_site x(S.lattice());
    forallsites(x) res+=S(x);
    mdp.add(res);
    return res/(S.lattice().nvol_gl);
}

```



```

};

int main(int argc, char **argv) {
    mdp.open_wormholes(argc, argv);

    int conf,Nconfig=10;
    int mybox[]={128, 128};
    mdp_lattice mylattice(2, mybox);
    ising_field      S(mylattice); /* ths spin field +1 or -1 */
    scalar_field     H(mylattice); /* the external magnetic field */
    mdp_site         x(mylattice);
    mdp_jackboot      jb(Nconfig,1);
    jb.plain(0);

    S.beta=0.5;          /* inverse square temperature */
    S.kappa=0.1;         /* coupling with the total external field */
    S.magnetic_field=1; /* an extra external magnetic field */

    for(S.beta=0.01; S.beta<0.15; S.beta+=0.001) {
        forallsites(x) {
            S(x)=1;
            H(x)=0;
        };
        S.update();
        H.update();

        for(conf=0; conf<Nconfig; conf++) {
            montecarlo_multihit(S,H);
            jb(conf,0)=average_spin(S);
        }
        mdp << "beta = " << S.beta << ", "
        << "average plaquette = " << jb.mean()
        << "(" << jb.j_err() << ")\n";
    }
    mdp.close_wormholes();
    return 0;
}

```




Figure 12: Total average magnetization in a two dimension spin system as function of the inverse temperature. The jump corresponds to a phase transition.

Note that h_{tot} is a constant therefore it is implemented as a member variable of **S** called **magnetic_field**.

The output of this program has been plotted and reported in fig. 12. It clearly shows that when the temperature is high ($\beta = 1/T$ is low) the average magnetization is insensible to the external magnetic field, while when the temperature is low (β is high) all the spins are oriented in the opposite to that of the magnetic field. The temperature where the phase transition (the jump) occurs is called the Debye temperature.

Note that the only reference to the number of dimensions is in the function **main()** in the lines

```
int mybox[]={128, 128};
mdp_lattice mylattice(2, mybox);
```

These are the only lines to change to generalize the program to more than 2D.

5.4 Solid under stress

Many other parallel applications could be developed relatively easy using MDP . For example one may be interested in studying the deformation of a solid under stress. The solid can be approximated with a finite lattice of given arbitrary shape and each site would be connected to its next-neighbor and next-next-neighbor sites with springs. One needs at least 4 fields on this lattice: M,K,G,Q and P, where M is the mass density of the solid; K the spring constant; G a parameter determining the internal friction of the solid; Q the physical (spatial) coordinates of each site, P the momentum associated to each site. The problem of determining the deformations of this solid can be solved by iterating a discretized form of the Hamilton equations

$$P_{t+\Delta t}(x) - P_t(x) = \frac{\partial H(P, Q, K, G)}{\partial Q(x)} \Delta t = F(x, Q, P, K, G) \Delta t \quad (14)$$

$$Q_{t+\Delta t}(x) - Q_t(x) = \frac{\partial H(P, Q, K, G)}{\partial P(x)} \Delta t = P(x)/M(x) \Delta t \quad (15)$$

These equations translates into

```
float delta_t=0.000001; /* a small nonzero positive number */
forallsites(x)
  for(mu=0; mu<3; mu+) {
    F[mu]=/* some function of K,G,Q,P */
    P(x,mu)=P(x,mu)+F(mu)*delta_t;
    Q(x,mu)=Q(x,mu)+P(x,mu)/m*delta_t;
  };
P.update();
Q.update();
```

The expression of the function F can be quite complicated depending on how one models the solid.

5.5 Lattice QCD

MDP has been used to develop a parallel package for large scale Lattice QCD simulations called **FermiQCD** [9, 10]. Here we only list some of its main features.

QCD	FermiQCD
Algebra of Euclidean gamma matrices	
$A = \gamma^\mu \gamma^5 e^{3i\gamma^2}$	<pre>mdp_matrix A; A=Gamma[mu]*Gamma5*exp(3*I*Gamma[2]);</pre>
Multiplication of a fermionic field for a spin structure	
$\forall x : \chi(x) = (\gamma^3 + m)\psi(x + \hat{\mu})$	<pre>/* assuming the following definitions mdp_lattice space_time(...); fermi_field chi(space_time,Nc); fermi_field psi(space_time,Nc); mdp_site x(space_time); */ forallsites(x) chi(x)=(Gamma[3]+m)*psi(x+mu);</pre>
Translation of a fermionic field	
$\forall x, a : \chi_a(x) = U(x, \mu)\psi_a(x + \hat{\mu})$	<pre>forallsites(x) for(a=0; a<psi.Nspin; a++) chi(x,a)=U(x,mu)*psi(x+mu,a);</pre>

Table 2: Example of FermiQCD statements.

The typical problem in QCD (Quantum Chromo Dynamics) is that of determining the correlation functions of the theory as a function of the parameters. From the knowledge of these correlation functions one can extract particle masses and matrix elements to compare with experimental results from particle accelerators.

On the lattice, each correlation function is computed numerically as the average of the corresponding operator applied to elements of a Markov chain of gauge field configurations. Both the processes of building the Markov chain and of measuring operators involve quasi-local algorithms. Some of the main features implemented in **FermiQCD** are:

- ☒ works on a single process PC
- ☒ works in parallel with MPI
- ☒ arbitrary lattice partitioning
- ☒ parallel I/O (partitioning independent)
- ☒ arbitrary space-time dimension
- ☒ arbitrary gauge group $SU(n)$
- ☒ anisotropic Wilson gauge action
- ☒ anisotropic Wilson fermionic action (optimized for P4)
- ☒ anisotropic Clover improved action (optimized for P4)
- ☒ domail-wall action
- ☒ Kogut-Susskind improved action (optimized for P4)
- ☒ ordinary and stochastic propagators
- ☒ minimal residue inversion
- ☒ stabilized biconjugate gradient inversion
- ☒ reads existing CANOPY/UKQCD/MILC data

In table 2 we show a few examples of FermiQCD Object Oriented capabilities (compared with examples in the standard textbook notation for Lattice QCD [10])




Figure 13: Efficiency e_1 as function of N and L , for $T = 60$ (the number 60 has been chosen because it has the biggest number of integer factors within the number of available processors).

6 Timing and efficiency issues

The issue of efficiency of a parallel program is not a simple one, in fact there are many factors to take into account. In particular the efficiency scales with the number of parallel processes, with the size of the data that need to be communicated between processes, and with the total volume of the lattice.

The dependence of the efficiency on these variables cannot be in general factored because of non-linearities introduced by modern computer architectures (for example caches effects, finite memory, memory speed/cpu speed, latencies, etc).

To give an idea of how a parallel computer program based on MDP scales with the number of processes, we consider the following program.

```
// program: application4.cpp
#include "mdp.h"
```




Figure 14: Efficiency e_2 as function of N , for $T = 2N$ and $L = 8$.

```

int main(int argc, char **argv) {
    mdp.open_wormholes(argc, argv);
    int box[4]={8,12,12,12};
    mdp_lattice space(4,box);
    mdp_matrix_field U(space,3,3);
    mdp_matrix_field V(space,3,3);
    mdp_site x(space);
    forallsites(x) {
        V(x)=space.random(x).SU(3);
        U(x)=0;
    }
    U.update();
    V.update();
    for(int i=0; i<10000; i++) {
        forallsites(x)
            U(x)=0.125*(cos(U(x)+V(x))+
                        U(x+0)+U(x-0)+

```

```

        U(x+1)+U(x-1)+
        U(x+2)+U(x-2)+
        U(x+3)+U(x-3));
/* uncomment to print convergence
if(onwhichprocess(space,0,0,0,0)==ME) {
    x.set(0,0,0,0);
    cout << real(U(x,0,0)) << endl;
}
*/
U.update();
}
V.save("V_field.dat");
U.save("U_field.dat");
mdp.close_wormholes();
return 0;
}

```

It solves a 4-dimensional Laplace equation for a field of 3×3 matrices, and performs some parallel IO as well. The reader can work out the details of the program as an exercise.

We ran the same program for different lattice volumes $V = T \times L^3$ and on different sets of processors (N =number of processors). Each processor (node) is a 700MHz PentiumIII PC. The nodes are connected using Myrinet cards/switches.

For each run we measured the total time $t(T, L, N)$ and we then computed the efficiency according with two possible definitions¹²:

$$e_1(T, L) = 100 \frac{t(T, L, 1)}{N t(T, L, N)} \quad (16)$$

$$e_2(N, L, x) = 100 \frac{t(x, L, 1)}{t(Nx, L, N)} \quad (17)$$

The first definition is applicable to computations in which the problem size is fixed and one varies the numbers of processors. The second definition is applicable to computations for which the size of the problem increases with the

¹²In the context of Lattice QCD simulation, up to an overall normalization, the first definition is the one used at Fermilab to benchmark parallel software, the second is the definition used by the MILC collaboration.

number of processors available. In both cases the efficiency is normalized to 100%. In an ideal world (without latencies and with instantaneous communications) both these efficiencies should be constant and equal to 100%.

In figures 13-14 we plot e_1 and e_2 for our runs and the maximum time spent in communications by the nodes (as computed by **MDP**) that we call r . The normalizations are given by

$$\begin{aligned} t(60, 4, 1) &= 788\text{sec.} \\ t(60, 8, 1) &= 6417\text{sec.} \\ t(60, 12, 1) &= 6417\text{sec.} \\ t(2, 8, 1) &= 201\text{sec.} \\ t(8, 8, 1) &= 834\text{sec.} \end{aligned}$$

These numbers can be further reduced by writing a user defined function that optimizes the critical computation done in the line

$$U(x) = 0.125 * (\cos(U(x) + V(x)) + \dots$$

From the plots we observe that with good approximation

$$e_i \simeq 100 - r \tag{18}$$

for both the definitions of efficiency.

Eventually, for N bigger than some N_0 , the total time becomes totally dominated by communication (the time spent in communications goes to 100%) and e_1 starts to scale as $1/N$.

Note that the answer to the question “how many processors should I use?” is not addressed in this paper because, in financial terms, this depends on the subjective utility function of each individual user. The only general conclusion is that N should be less than N_0 .

A different algorithm may scale in a different way even if the qualitative behavior of the efficiency should be the same. Some algorithms may require more communications than others and may result in a bigger loss of efficiency when running in parallel.

Acknowledgements

I thank Chris Sachrajda, Jonathan Flynn and Luigi Del Debbio of the University of Southampton (UK) where this project started.

The development of the code `MDP` in its final (but not yet definitive) form and its main application, `FermiQCD`, have greatly benefit from discussions and suggestions with members of the Fermilab Theory Group. In particular I want to thank Paul Mackenzie, Jim Simone, Jimmy Juge, Andreas Kronfeld and Estia Eichten for comments and suggestions. Moreover I freely borrowed many ideas from existing Lattice QCD codes (including the UKQCD code [11], the MILC code [12], QCDF90 [13] and CANOPY [14]). I here thank their authors for making their software available to me.

A Syntax of MDP

Only the most important functions.

```
#ifndef USE_DOUBLE_PRECISION
typedef mdp_real float;
#else
typedef mdp_real double;
#endif

#define and      &&
#define or      ||
#define Pi      3.14159265359
#define I      mdp_complex(0,1)
#define CHECK_ALL
#define TRUE  1
#define FALSE 0
#define NOWHERE 268435456
#define EVEN  0
#define ODD   1
#define EVENODD 2

#define ME mdp.id()
#define Nproc mdp.nproc()

#define forallsites(x)
#define forallsitesofparity(x,parity)
#define forallsitesandcopies(x)
```

```

#define forallsitesandcopiesofparity(x,parity)

class mdp_complex {
    mdp_real re, im;
};

template<class T, int ndim>
class mdp_array {
public:
    mdp_array()
    mdp_array(int d[])
    mdp_array(int d0, int d1, ...)
    long size()
    int size(int mu);
    T* address();
    dimension(int d[]);
    dimension(int d0, int d1, ...);
    T &operator() (int i, int j, ...);
};

class mdp_matrix {
public:
    mdp_matrix();
    mdp_matrix(int r, int c);
    long size();
    int rowmax();
    int colmax();
    mdp_complex *address();
    int dimension(int rows, int cols);
    mdp_complex &operator() (int row, int col);
    mdp_matrix operator() (int row);
    // operator+;
    // operator-;
    // operator*;
    // operator/;
    // and various functions ....
};

```

```

class mdp_jackboot {
    void *handle;
    float (*f)(float*,void*)
    float &operator() (int config, int operator);
    float plain(int i);
    float mean();
    float j_err();
    float b_err(int bootsamples=200);
};

class mdp_random_generator {
public:
    float plain();
    float gaussssian();
    float distribution(float (*f)(float,void*),void*);
    mdp_matrix SU(int n);
};

class mdp_communicator {
open_wormholes(int&, char***);
close_wormholes();
template <class T> void put(T& object, int destination);
template <class T> void put(T& object, int destination,
                           mpi_request&);
template <class T> void get(T& object, int source);
template <class T> void put(T* object, long size, int destination);
template <class T> void put(T* object, long size, int destination,
                           mpi_request&);
template <class T> void get(T* object, long size, int source);
void wait(mpi_request&);
void wait(mpi_request*, int size);
void add(float& source_obj, float& dest_obj);
void add(float* source_obj, float* dest_obj, long size);
void add(double& source_obj, double& dest_obj);
void add(double* source_obj, double* dest_obj, long size);
void add(int&);

```

```

void add(long&);
void add(float&);
void add(double&);
void add(mdp_complex&);
void add(mdp_matrix&);
void add(int&,long);
void add(long&,long);
void add(float&,long);
void add(double&,long);
void add(mdp_complex&,long);
void add(mdp_matrix&,long);
template <class T> void broadcast(T& object, int source);
template <class T> void broadcast(T* object, long size, int source);
void barrier();
double time();
void reset_time();
wait abort();
} mpi;

class mdp_lattice {
int ndim;
long nvol;
mdp_lattice(int ndim,
            int *nx,
            int (*where)(int*,int,int*)=default_partitioning<0>,
            void (*topology)(int,int*,int*,int*,int,int*)=torus_topology,
            long seed=0,
            int next_next=1,
            bool local_random_);
mdp_lattice(int ndim,
            int ndir,
            int *nx,
            int (*where)(int*,int,int*)=default_partitioning<0>,
            void (*topology)(int,int*,int*,int*,int,int*)=torus_topology,
            long seed=0,
            int next_next=1,
            bool local_random_);

```

```

void allocate_lattice(int ndim,
    int *nx,
    int (*where)(int*,int,int*)=default_partitioning<0>,
    void (*topology)(int,int*,int*,int*,int,int*)=torus_topology,
    long seed=0,
    int next_next=1,
    bool local_random_);
void allocate_lattice(int ndim,
    int ndir,
    int *nx,
    int (*where)(int*,int,int*)=default_partitioning<0>,
    void (*topology)(int,int*,int*,int*,int,int*)=torus_topology,
    long seed=0,
    int next_next=1,
    bool local_random_);
long global_coordinate(int *);
void global_coordinate(long, int *);
int compute_parity(int*);
void deallocate_memory();
void initialize_random(long seed);
mdp_random_generator& random(site x);
long local_volume();
long global_volume();
int n_dimensions();
int n_directions();
long size();
long size(int);
long local_volume();
long global_volume();
long move_up(long, int);
long move_down(long,int);
long start_index(int, int);
long stop_index(int,int);
};

int on_which_process(mdp_lattice &lattice,
    int x0, int x1, int x2, ...);

```

```

class vector {
    int x[10];
};

vector bynary2versor(long);
int    versor2binary(int,int,int,int,int...);
long   vector2binary(vector);

class mdp_site {
site(mdp_lattice& mylattice);
void start(int parity);
void next();
int is_in();
int is_here();
int is_in_boundary();
int is_equal(int x0, int x1, int x2, ...);
int parity();
long local_index();
long global_index();
void set(int x0, int x1, int x2, ...);
void set_local(long index);
void set_global(long index);
mdp_site operator+(int mu);
mdp_site operator-(int mu);
int operator()(int mu);
int operator=(mdp_site x);
int operator=(int *x);
int operator==(mdp_site x);
int operator==(int *x);
int operator!=(mdp_site x);
int operator!=(int *x);
int operator=(vector);
int operator+(vector);
int operator-(vector);
};

```

```

long site2binary(mdp_site);

template<class T>
class mdp_field {
mdp_field();
mdp_field(mdp_lattice& mylattice,
          int field_components);
void allocate_field(mdp_lattice& mylattice,
                   int field_components);
void deallocate_memory();
T* address(mdp_site x);
T* address(mdp_site x, int field_component);
void operator=(mdp_field& psi);
mdp_lattice& lattice();
void operator=(const mdp_field&);
void operator=(const T);
void operator+=(const mdp_field&);
void operator==(const mdp_field&);
void operator+=(const mdp_field&);
template<class T2>
void operator*=(const T2);
template<class T2>
void operator/=(const T2);
void update(int parity, int block_n, int block_size);
void load(char filename[], int processI0, long buffersize);
void save(char filename[], int processI0, long buffersize);
long physical_size();
long size_per_site();
long physical_local_start(int);
long physical_local_stop(int);
T* physical_address()
T& physical_address(long)
void assign_to_null();
};

```


References

- [1] G. Satir and D. Brown, C++ The Core Language, O'Reilly and Associates (1995)
- [2] B. Stroustup, C++ Programming Language, Addison-Welley (1997)
- [3] N. Cabibbo and E. Marinari, A new method for updating $SU(N)$ matrices in computer simulations of gauge theories, Phys. Lett. **119B** (1982) 387
- [4] J. Shao and D. Tu, The Jackknife and Bootstrap, Springer Verlag (1995)
- [5] P. D. Coddington, Random Number Generators for Parallel Computers (1997) (unpublished)
- [6] G. A. Marsaglia, A current view on random number generators, in *Computational Science and statistics: The Interface*, ed. L. Balliard, Elsevier, Amsterdam (1985)
- [7] P. S. Pacheco, Parallel Programming with MPI, San Francisco, CA, Morgan Kaufmann (1997)
- [8] G. Banhot, The Metropolis Algorithm, Rep. Prog. Phys. 51 (1988) 429
- [9] M. Di Pierro, "Matrix distributed processing and FermiQCD", hep-lat/0011083.
- [10] M. Di Pierro, "From Monte Carlo integration to lattice quantum chromodynamics: An introduction," hep-lat/0009001
- [11] The UKQCD collaboration web page:
<http://www.ph.ed.ac.uk/ukqcd/>
- [12] The MILC collaboration web page:
<http://physics.indiana.edu/~sg/milc.html>
- [13] I. Dasgupta, A.R. Levi, V. Lubicz and C. Rebbi, Comput. Phys. Commun. 98 (1996) 365-397

- [14] M. Fischler, J. Hockney, P. Mackenzie and M. Uchima, “Canopy 7.0 manual”, Fermilab preprint TM-1881. PDF text available at web site:
<http://fnalpubs.fnal.gov/archive/tm/TM-1881.pdf>